

1. Download all the data in this folder <https://drive.google.com/open?id=1Z4TyI7FcFVEx8qdl4jO9qxvxaqLSqoEu>. it contains two file both images and labels. The label file list the images and their categories in the following format:

**path/to/the/image.tif,category**

where the categories are numbered 0 to 15, in the following order:

- 0 letter
- 1 form
- 2 email
- 3 handwritten
- 4 advertisement
- 5 scientific report
- 6 scientific publication
- 7 specification
- 8 file folder
- 9 news article
- 10 budget
- 11 invoice
- 12 presentation
- 13 questionnaire
- 14 resume
- 15 memo

2. On this image data, you have to train 3 types of models as given below. You have to split the data into Train and Validation data.

3. Try not to load all the images into memory, use the generators that we have given the reference notebooks to load the batch of images only during the train data.

or you can use this method also

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-imagedatagenerator-with-flow-from-dataframe-8bd5776e45c1>

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c>

4. You are free to choose Learning rate, optimizer, loss function, image augmentation, any hyperparameters. but you have to use the same architecture what we are asking below.

5. Use tensorboard for every model and analyse your gradients. (you need to upload the screenshots for each model for evaluation)

Note: `fit_generator()` method will have problems with the tensorboard histograms, try to debug it, if you could not do use `histograms=0` i.e don't include histograms, check the documentation of tensorboard for more information.

6. You can check about Transfer Learning in this link - <https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>

## Model-1

## Model-2

## Model-3

In [1]:

```
import tensorflow as tf
import os
```

```
import numpy as np
import pandas as pd
```

In [2]:

```
from tensorflow.keras.layers import
Dense, Input, Conv2D, MaxPool2D, Activation, Dropout, Flatten, GlobalAveragePooling2D
from tensorflow.keras.models import Model
import random as rn
```

In [3]:

```
!gdown --id 1Z4TyI7FcFVEx8qdl4jO9qxvxaqLSqoEu
```

```
get_ipython().system_raw("unrar x rv1-cdip.rar")
```

Downloading...  
From: <https://drive.google.com/uc?id=1Z4TyI7FcFVEx8qdl4jO9qxvxaqLSqoEu>  
To: /content/rv1-cdip.rar  
4.66GB [01:25, 54.8MB/s]

In [4]:

```
os.listdir()
```

Out[4]:

```
['.config', 'data_final', 'rv1-cdip.rar', 'labels_final.csv', 'sample_data']
```

In [5]:

```
import pandas as pd

data = pd.read_csv('labels_final.csv') #reading the csv file
```

In [6]:

```
data.head()
```

Out[6]:

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	3
1	imagesl/l/x/t/xt19d00/502213303.tif	3
2	imagesx/x/e/d/xed05a00/2075325674.tif	2
3	imageso/o/j/b/ojb60d00/517511301+-1301.tif	3
4	imagesq/q/z/k/qzk17e00/2031320195.tif	7

In [7]:

```
from sklearn.model_selection import train_test_split

#traindf, validationdf = train_test_split(data, test_size=0.3, random_state=42, shuffle=True)

train_path, validation_path, train_label, validation_label = train_test_split(data['path'], data['label'], test_size=0.2, random_state=42)
```

In [8]:

```
labels_dict = { 0 : 'letter', 1: 'form', 2: 'email', 3 : 'handwritten', 4 : 'advertisement',
                5 : 'scientific report', 6 : 'scientific publication', 7 : 'specification', 8 : 'file folder',
                9 : 'news article', 10 : 'budget', 11 : 'invoice', 12 : 'presentation',
                13 : 'contract', 14 : 'receipt', 15 : 'memo' }
```

```
13 : 'questionnaire', 14 : 'resume', 15 : 'memo'}
```

In [9]:

```
labels_dict.values()
```

Out[9]:

```
dict_values(['letter', 'form', 'email', 'handwritten', 'advertisement', 'scientific report', 'scientific publication', 'specification', 'file folder', 'news article', 'budget', 'invoice', 'presentation', 'questionnaire', 'resume', 'memo'])
```

In [10]:

```
for subfolder_name in list(labels_dict.values()):  
    os.makedirs(os.path.join('train_images', subfolder_name))
```

In [11]:

```
os.listdir('train_images')
```

Out[11]:

```
['advertisement',  
 'handwritten',  
 'memo',  
 'form',  
 'invoice',  
 'email',  
 'file folder',  
 'scientific publication',  
 'presentation',  
 'budget',  
 'news article',  
 'questionnaire',  
 'letter',  
 'resume',  
 'scientific report',  
 'specification']
```

In [12]:

```
#https://thispointer.com/python-how-to-copy-files-from-one-location-to-another-using-shutil-copy/  
import shutil  
from tqdm import tqdm  
for file, label in tqdm(zip(train_path, train_label)):  
    shutil.copy('data_final/'+file, 'train_images/'+labels_dict[label]+'/')
```

```
38400it [02:08, 299.08it/s]
```

In [13]:

```
for subfolder_name in list(labels_dict.values()):  
    os.makedirs(os.path.join('validation_images', subfolder_name))
```

In [14]:

```
for file, label in tqdm(zip(validation_path, validation_label)):  
    shutil.copy('data_final/'+file, 'validation_images/'+labels_dict[label]+'/')
```

```
9600it [00:34, 279.55it/s]
```

In [15]:

```
dir_path= 'train_images'  
for i in os.listdir(dir_path):  
    print("No of Images in ",i," category is ",len(os.listdir(os.path.join(dir_path,i))))
```

```
No of Images in advertisement category is 2393
No of Images in handwritten category is 2413
No of Images in memo category is 2401
No of Images in form category is 2407
No of Images in invoice category is 2397
No of Images in email category is 2379
No of Images in file folder category is 2351
No of Images in scientific publication category is 2373
No of Images in presentation category is 2428
No of Images in budget category is 2422
No of Images in news article category is 2392
No of Images in questionnaire category is 2395
No of Images in letter category is 2461
No of Images in resume category is 2415
No of Images in scientific report category is 2379
No of Images in specification category is 2391
```

In [16]:

```
dir_path= 'validation_images'
for i in os.listdir(dir_path):
    print("No of Images in ",i," category is ",len(os.listdir(os.path.join(dir_path,i))))
```

```
No of Images in advertisement category is 601
No of Images in handwritten category is 592
No of Images in memo category is 595
No of Images in form category is 587
No of Images in invoice category is 595
No of Images in email category is 614
No of Images in file folder category is 652
No of Images in scientific publication category is 612
No of Images in presentation category is 578
No of Images in budget category is 580
No of Images in news article category is 609
No of Images in questionnaire category is 612
No of Images in letter category is 554
No of Images in resume category is 590
No of Images in scientific report category is 620
No of Images in specification category is 609
```

In [17]:

```
dir_path= 'train_images'
ImageFlow = tf.keras.preprocessing.image.ImageDataGenerator()
ImageGenerator_train =
ImageFlow.flow_from_directory(dir_path,target_size=(224,224),seed=10,batch_size=32)
```

Found 38397 images belonging to 16 classes.

In [18]:

```
dir_path='validation_images'
ImageGenerator_validation =
ImageFlow.flow_from_directory(dir_path,target_size=(224,224),seed=10,batch_size=32)
```

Found 9600 images belonging to 16 classes.

In [19]:

```
##Checking time taken to load images.
import time
start = time.time()
total_batches = 0

batches = 0
per_batch = 32
for x_batch, y_batch in ImageGenerator_train:
    batches += 1
    if batches >= 6899/per_batch:
        total_batches = total_batches + batches
        break
```

```

end = time.time()
duration = end-start
print("{} batches: {} s".format(total_batches, duration))
print("{:0.5f} Images/s".format(per_batch*total_batches/duration))

```

216 batches: 34.13801717758179 s  
202.47222 Images/s

1. Use [VGG-16](#) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, add a new Conv block ( 1 Conv layer and 1 Maxpooling ), 2 FC layers and a output layer to classify 16 classes. You are free to choose any hyperparameters/parameters of conv block, FC layers, output layer.
3. Final architecture will be **INPUT --> VGG-16 without Top layers(FC) --> Conv Layer --> Max pool Layer --> 2 FC layers --> Output Layer**
4. Train only new Conv block, FC layers, output layer. Don't train the VGG-16 network.



In [20]:

```

import os
os.environ['PYTHONHASHSEED'] = '0'

##https://keras.io/getting-started/faq/#how-can-i-obtain-reproducible-results-using-keras-during-d
## Have to clear the session. If you are not clearing, Graph will create again and again and graph
## Variables will also set to some value from before session
tf.keras.backend.clear_session()

## Set the random seed values to regenerate the model.
np.random.seed(0)
rn.seed(0)

#VGG-16
base_model=tf.keras.applications.VGG16(weights='imagenet', include_top=False,input_shape=(224,224,3
))

# add a global spatial average pooling layer
x=base_model.output

#print(x)

# first: train only the top layers (which were randomly initialized)
# i.e. freeze all convolutional InceptionV3 layers
for layer in base_model.layers:
    layer.trainable = False

#base_model.trainable=False

#Conv Layer
Conv1 = Conv2D(filters=32,kernel_size=(3,3),strides=(1,1),padding='valid',data_format='channels_las
t',
               activation='relu',kernel_initializer=tf.keras.initializers.he_normal(seed=0),name='C
onv1')(x)
#MaxPool Layer
Pool1 = MaxPool2D(pool_size=(2,2),strides=(2,2),padding='valid',data_format='channels_last',name='P
ool1')(Conv1)

flatten = Flatten(data_format='channels_last',name='Flatten')(Pool1)

#FC layer
FC1 = Dense(units=64,activation='relu',kernel_initializer=tf.keras.initializers.glorot_normal(seed=
32),name='FC1')(flatten)

#FC layer
FC2 = Dense(units=32,activation='relu',kernel_initializer=tf.keras.initializers.glorot_normal(seed=
33),name='FC2')(FC1)

#output layer

```

```
#Output layer
Out =
Dense(units=16,activation='softmax',kernel_initializer=tf.keras.initializers.glorot_normal(seed=3)
,name='Output')(FC2)

#Creating a model
model_1 = Model(inputs=base_model.input,outputs=Out)
```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)  
58892288/58889256 [=====] - 0s 0us/step

In [21]:

```
model_1.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[ (None, 224, 224, 3) ]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
Conv1 (Conv2D)	(None, 5, 5, 32)	147488
Pool1 (MaxPooling2D)	(None, 2, 2, 32)	0
Flatten (Flatten)	(None, 128)	0
FC1 (Dense)	(None, 64)	8256
FC2 (Dense)	(None, 32)	2080
Output (Dense)	(None, 16)	528
=====		
Total params: 14,873,040		
Trainable params: 158,352		
Non-trainable params: 14,714,688		

In [22]:

```
model_1.compile(optimizer='rmsprop', loss='categorical_crossentropy',metrics=["accuracy"])
```

In [23]:

```
history=model_1.fit(ImageGenerator_train,  
  
validation_data=ImageGenerator_validation,epochs=3,steps_per_epoch=len(ImageGenerator_train),validation_steps=len(ImageGenerator_validation))
```

```
Epoch 1/3  
1200/1200 [=====] - 277s 203ms/step - loss: 2.2030 - accuracy: 0.3689 - val_loss: 1.5048 - val_accuracy: 0.5334  
Epoch 2/3  
1200/1200 [=====] - 222s 185ms/step - loss: 1.4179 - accuracy: 0.5628 - val_loss: 1.4221 - val_accuracy: 0.5724  
Epoch 3/3  
1200/1200 [=====] - 188s 157ms/step - loss: 1.2840 - accuracy: 0.6057 - val_loss: 1.4315 - val_accuracy: 0.5847
```

In [24]:

```
history.history
```

Out[24]:

```
{'accuracy': [0.45391565561294556, 0.5649399757385254, 0.6053076982498169],  
'loss': [1.7981104850769043, 1.4176541566848755, 1.3013049364089966],  
'val_accuracy': [0.5334374904632568, 0.5723958611488342, 0.5846874713897705],  
'val_loss': [1.5048243999481201, 1.422074317932129, 1.4315128326416016]}
```

In [25]:

```
from keras.preprocessing.image import ImageDataGenerator
```

In [ ]:

```
df=pd.read_csv('./labels_final.csv',dtype=str)  
traindf, validationdf = train_test_split(df, test_size=0.3, random_state=42,shuffle=True)
```

```
Found 0 validated image filenames belonging to 0 classes.  
Found 0 validated image filenames belonging to 0 classes.
```

```
/usr/local/lib/python3.7/dist-packages/keras_preprocessing/image/dataframe_iterator.py:282:  
UserWarning: Found 33600 invalid image filename(s) in x_col="path". These filename(s) will be ignored.  
    .format(n_invalid, x_col)  
/usr/local/lib/python3.7/dist-packages/keras_preprocessing/image/dataframe_iterator.py:282:  
UserWarning: Found 33600 invalid image filename(s) in x_col="path". These filename(s) will be ignored.  
    .format(n_invalid, x_col)
```

In [ ]:

```
datagen=ImageDataGenerator(rescale=1./255.)  
  
train_generator=datagen.flow_from_dataframe(dataframe=traindf,directory='./train_images/',x_col="path",  
y_col="label",batch_size=32,seed=42,shuffle=False,class_mode="categorical",target_size=(224,224,3),  
subset='training')  
  
valid_generator=datagen.flow_from_dataframe(dataframe=traindf,directory='./validation_images/',x_col="path",  
y_col="label",batch_size=32,seed=42,shuffle=False,class_mode="categorical",target_size=(224,224,3),  
subset='validation')  
  
# Define your model
```

```
# model.compile(optimizer='adam', loss=....., metrics=["accuracy"])

#
model.fit(train_generator, validation_data=validation_generator, epochs=10, steps_per_epoch=len(train_generator), validation_steps=len(validation_generator))
```

1. Use [VGG-16](#) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, don't use FC layers, use conv layers only as Fully connected layer. any FC layer can be converted to a CONV layer. This conversion will reduce the No of Trainable parameters in FC layers. For example, an FC layer with K=4096 that is looking at some input volume of size 7×7×512 can be equivalently expressed as a CONV layer with F=7, P=0, S=1, K=4096. In other words, we are setting the filter size to be exactly the size of the input volume, and hence the output will simply be 1×1×4096 since only a single depth column "fits" across the input volume, giving identical result as the initial FC layer. You can refer [this](#) link to better understanding of using Conv layer in place of fully connected layers.
3. Final architecture will be VGG-16 without FC layers(without top), 2 Conv layers identical to FC layers, 1 output layer for 16 class classification. **INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer**
3. Train only last 2 Conv layers identical to FC layers, 1 output layer. Don't train the VGG-16 network.

In [39]:

```
import os
os.environ['PYTHONHASHSEED'] = '0'

tf.keras.backend.clear_session()

## Set the random seed values to regenerate the model.
np.random.seed(0)
rn.seed(0)

#VGG-16
base_model_2=tf.keras.applications.VGG16(weights='imagenet', include_top=False, input_shape=(224,224,3))
print(base_model_2.input)

# add a global spatial average pooling layer
x=base_model_2.output

for layer in base_model_2.layers:
    layer.trainable = False

Conv1 = Conv2D(4096, kernel_size=[7,7], strides=(1,1), padding='valid', activation='relu', name='Conv1')(x)

Conv2 = Conv2D(4096, kernel_size=[1,1], strides=(1,1), padding='valid', data_format='channels_last', activation='relu', kernel_initializer=tf.keras.initializers.he_normal(seed=0), name='Conv2')(Conv1)

flatten = Flatten(data_format='channels_last', name='Flatten')(Conv2)
output = Dense(units=16, activation='softmax', kernel_initializer=tf.keras.initializers.glorot_normal(seed=3), name='Output')(flatten)

# #Creating a model
model_2 = Model(inputs=base_model_2.input, outputs=output)
```

```
KerasTensor(type_spec=TensorSpec(shape=(None, 224, 224, 3), dtype=tf.float32, name='input_1'), name='input_1', description="created by layer 'input_1'")
```

In [40]:



```
model_2.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
Conv1 (Conv2D)	(None, 1, 1, 4096)	102764544
Conv2 (Conv2D)	(None, 1, 1, 4096)	16781312
Flatten (Flatten)	(None, 4096)	0
Output (Dense)	(None, 16)	65552
=====		
Total params: 134,326,096		
Trainable params: 119,611,408		
Non-trainable params: 14,714,688		

In [41]:

```
model_2.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])

history_2=model_2.fit(ImageGenerator_train,

validation_data=ImageGenerator_validation,epochs=3,steps_per_epoch=len(ImageGenerator_train),validation_steps=len(ImageGenerator_validation))
```

Epoch 1/3

1200/1200 [=====] - 209s 173ms/step - loss: 21.8210 - accuracy: 0.3893 - val\_loss: 1.4671 - val\_accuracy: 0.6195

Epoch 2/3

1200/1200 [=====] - 207s 173ms/step - loss: 1.2502 - accuracy: 0.6540 - val\_loss: 1.2839 - val\_accuracy: 0.6566

Epoch 3/3

1200/1200 [=====] - 207s 173ms/step - loss: 1.0990 - accuracy: 0.7062 - v

al\_loss: 1.2574 - val\_accuracy: 0.6760

1. Use same network as Model-2 'INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer' and train only Last 6 Layers of VGG-16 network, 2 Conv layers identical to FC layers, 1 output layer.

### MODEL 3

In [43]:

```
import os
os.environ['PYTHONHASHSEED'] = '0'

tf.keras.backend.clear_session()

## Set the random seed values to regenerate the model.
np.random.seed(0)
rn.seed(0)

#VGG-16
base_model_3=tf.keras.applications.VGG16(weights='imagenet', include_top=False,input_shape=(224,224,3))

# add a global spatial average pooling layer
x=base_model_3.output

for layer in base_model_3.layers[:13]:
    layer.trainable = False

## #Conv Layer
Conv1 = Conv2D(4096, kernel_size=[7,7], strides=(1,1), padding='valid', activation='relu', name='Conv1')(x)

Conv2 = Conv2D(4096, kernel_size=[1,1], strides=(1,1), padding='valid', data_format='channels_last', activation='relu', kernel_initializer=tf.keras.initializers.he_normal(seed=0), name='Conv2')(Conv1)
flatten = Flatten(data_format='channels_last', name='Flatten')(Conv2)
output = Dense(units=16, activation='softmax', kernel_initializer=tf.keras.initializers.glorot_normal(seed=3), name='Output')(flatten)

## Creating a model
model_3 = Model(inputs=base_model_3.input, outputs=output)
```

In [44]:

```
model_3.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[ (None, 224, 224, 3) ]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0

block3_pool (MaxPooling2D)	(None, 56, 56, 256)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
Conv1 (Conv2D)	(None, 1, 1, 4096)	102764544
Conv2 (Conv2D)	(None, 1, 1, 4096)	16781312
Flatten (Flatten)	(None, 4096)	0
Output (Dense)	(None, 16)	65552
=====		
Total params: 134,326,096		
Trainable params: 129,050,640		
Non-trainable params: 5,275,456		

In [45]:

```
model_3.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])
history_3=model_3.fit(ImageGenerator_train,
validation_data=ImageGenerator_validation, epochs=3, steps_per_epoch=len(ImageGenerator_train), validation_steps=len(ImageGenerator_validation))
```

```
Epoch 1/3
1200/1200 [=====] - 284s 236ms/step - loss: 227.2454 - accuracy: 0.0603 - val_loss: 2.7729 - val_accuracy: 0.0604
Epoch 2/3
1200/1200 [=====] - 284s 236ms/step - loss: 2.7727 - accuracy: 0.0593 - val_loss: 2.7731 - val_accuracy: 0.0577
Epoch 3/3
1200/1200 [=====] - 284s 236ms/step - loss: 2.7728 - accuracy: 0.0641 - val_loss: 2.7731 - val_accuracy: 0.0577
```

In [46]:

```
history_3.history
```

Out[46]:

```
{'accuracy': [0.06013490632176399, 0.0614110492169857, 0.06370289623737335],
 'loss': [39.831398010253906, 2.772803544998169, 2.7727930545806885],
 'val_accuracy': [0.06041666492819786,
 0.05770833417773247,
 0.05770833417773247],
 'val_loss': [2.772860527038574, 2.773062229156494, 2.7731451988220215]}
```

## MODEL 1 Accuracy

In [50]:

```
history.history['accuracy'][2]*100
```

Out[50]:

60.53076982498169

#### MODEL 2 accuracy

In [51]:

```
history_2.history['accuracy'][2]*100
```

Out[51]:

70.18256783485413

#### MODEL 3 accuracy

In [52]:

```
history_3.history['accuracy'][2]*100
```

Out[52]:

6.370289623737335

The Last model Accuracy was low because we tried to train the weight of the vgg model with our own data and thus it didnt perform well.